

Who Is Responsible When AI Makes a Decision?

GRADES
9-12

Grades 9-12 | 55-75 minutes | Evidence, fairness and human judgment

AI may recommend. People remain accountable.

Investigate a fictional internship recommendation, uncover missing context, and design a process that people can review, explain and correct.



ESSENTIAL QUESTION

When AI influences a high-impact decision, who is responsible for the outcome - and what safeguards should be required?

Classroom lesson package | Includes educator guidance, student tools and assessment

TEACHER GUIDE

Lesson at a Glance

A complete, adaptable lesson for high school classrooms

LESSON LENGTH

55-75 minutes

GROUP STRUCTURE

Teams of 3-5 students

MATERIALS

Case study, evidence cards, stakeholder cards, decision framework, rubric, exit ticket

TECHNOLOGY

Optional; lesson can be completed without student AI accounts

PREPARATION

Print or distribute student materials; select discussion protocol

ASSESSMENT

Team recommendation, written rationale, exit ticket

Learning Objectives

- Explain how an AI-supported system can influence a school decision without independently making the final decision.
- Identify data limitations, missing context, uneven weighting, and potential fairness concerns.
- Distinguish between an AI recommendation and accountable human decision-making.
- Evaluate a decision process from multiple stakeholder perspectives.
- Design safeguards that improve transparency, appeal rights, human review, and accountability.
- Defend a recommendation using evidence and ethical reasoning.

Student-Friendly Success Criteria

- I can explain what the system was asked to do and what information it used.
- I can identify information that may be missing, incomplete, or weighted unfairly.
- I can explain who could be affected by the decision.
- I can propose a process in which people review, explain, and remain accountable for the outcome.
- I can support my recommendation with evidence from the case.

FACILITATOR NOTE

This lesson does not require students to decide whether AI is "good" or "bad." The instructional goal is to evaluate the quality and accountability of a specific decision process.

TEACHER GUIDE

Facilitation Plan

A 60-minute core sequence; extend discussion to 75 minutes as needed

Phase	Minutes	Facilitation move
Launch	5	Frame the question; distinguish recommendation from final decision.
Read the scenario	5	Read the case aloud or individually. No video is required.
Notice and wonder	4	Record observations, questions and missing information.
Investigate evidence	14	Analyze all ten evidence cards using the five-part framework.
Consider stakeholders	8	Compare perspectives and likely effects.
Design safeguards	12	Name reviewers, decision owner, correction pathway and pause trigger.
Share and challenge	8	Cite evidence, respond to a counter-question and revise.
Exit ticket	4	Explain human accountability and a useful safeguard.

OPENING SCRIPT

"Today we are asking whether the full decision process is responsible, fair, explainable and accountable."

DISCUSSION NORMS

Critique the process; separate evidence from assumptions; acknowledge uncertainty; ask who can challenge the result. A numerical output is not automatically neutral.

STUDENT CASE STUDY

The Internship Recommendation

A fictional scenario for evidence-based analysis

The Situation

Riverview High School created a competitive summer internship program with local employers. The program has 40 available placements and more than 150 applicants. To help staff review applications, the school begins using an AI-supported recommendation system. The system does not officially select students, but it assigns each applicant a recommendation level that staff members see before making placement decisions.

WHAT THE SYSTEM REVIEWS

- Course history and grades
- Attendance records
- Teacher recommendations
- Extracurricular and service activities
- Student career-interest survey responses
- Prior participation in school programs

The Unexpected Result

Jordan, an eleventh-grade student, applies for a technology and operations internship. Jordan has completed a computer science course, helps repair devices for family members, works part-time, and has strong interest in technology careers. Jordan is not listed among the system's recommended candidates.

A counselor reviews the result and notices several details: Jordan's part-time work was entered only as "employment," the application did not include a field for informal technical experience, and one teacher recommendation was submitted after the system processed the initial application data. Jordan also had several absences related to a documented family responsibility.

YOUR TASK

Team Challenge

Determine whether the school should rely on the recommendation as it currently stands. Then design a responsible decision process that uses technology appropriately while protecting students from incomplete, unfair, or unexplained outcomes.

STUDENT MATERIALS

Evidence Cards

Cut apart or distribute digitally; teams should examine every card

EVIDENCE 1 - COURSEWORK

Jordan completed Introductory Computer Science with a B and earned strong project scores. The system gave higher weight to advanced technical courses.

EVIDENCE 2 - ATTENDANCE

Jordan had nine absences. Six were connected to documented family responsibilities. The system used total absences without considering context.

EVIDENCE 3 - WORK EXPERIENCE

Jordan works 15 hours each week at a grocery store. The application listed the job but did not capture technology-related tasks Jordan performs outside school.

EVIDENCE 4 - INFORMAL SKILLS

Jordan repairs laptops and phones for relatives and helped a community organization set up devices. The application had no field for informal or community-based technical experience.

EVIDENCE 5 - RECOMMENDATION TIMING

A teacher submitted a strong recommendation after the first system review. Staff members are unsure whether the score was recalculated.

EVIDENCE 6 - CAREER INTEREST

Jordan selected information technology, operations, and entrepreneurship as top career interests.

EVIDENCE 7 - SYSTEM TRANSPARENCY

Counselors can see the recommendation category but cannot see how each factor influenced the result.

EVIDENCE 8 - HUMAN REVIEW

Program staff planned to review all applicants, but the recommended list is displayed first and may shape their attention.

EVIDENCE 9 - APPEAL PROCESS

The application materials do not explain how a student can correct information or challenge a recommendation.

EVIDENCE 10 - PROGRAM GOAL

The internship program is intended to identify students with readiness, interest, growth potential, and the ability to benefit from workplace learning.

STUDENT MATERIALS

Stakeholder Perspective Cards

Use one card per team member or rotate perspectives

JORDAN - STUDENT APPLICANT

Wants a fair review, a clear explanation, and an opportunity to correct missing information. Jordan is concerned that experiences outside formal courses were overlooked.

SCHOOL COUNSELOR

Wants students to receive equitable opportunities and needs enough information to explain recommendations. The counselor is worried about overreliance on the system.

INTERNSHIP PROGRAM DIRECTOR

Must review many applications efficiently while maintaining quality and fairness. The director wants a process that is practical and defensible.

TEACHER RECOMMENDER

Believes Jordan has strong problem-solving skills and persistence. The teacher wants late-arriving evidence to be considered.

DISTRICT TECHNOLOGY TEAM

Wants the system to be secure, reliable, and used only for an approved purpose. The team needs clear documentation and monitoring.

FAMILY MEMBER

Wants to understand why Jordan was not recommended and how the decision can be reviewed. The family is concerned about privacy and fairness.

EMPLOYER PARTNER

Wants students who can learn, collaborate, and contribute. The employer values potential and transferable skills, not only formal coursework.

DISTRICT LEADER

Must ensure that the process is educationally sound, legally and ethically responsible, and aligned with district values.

STUDENT TOOL

Responsible AI Decision Framework

Use the five questions to evaluate the full decision process

01 PURPOSE

What decision is supported? Is AI appropriate? What must it never decide alone?

.....

.....

02 DATA

Is information accurate, current, relevant and complete? What context is missing?

.....

.....

03 FAIRNESS

Who may benefit or be disadvantaged? Whose experience is overlooked?

.....

.....

04 TRANSPARENCY

Can people understand the result and correct errors?

.....

.....

05 HUMAN OVERSIGHT

Who reviews, decides, explains, and remains accountable?

.....

.....

Recommendation: accept, pause or reject the current use? Cite evidence card numbers.

.....

STUDENT TOOL

Design a Responsible Decision Process

Build safeguards into each stage of the process

Decision stage	Required safeguard	Accountable person / team
1 Before use	Define purpose, approved data, prohibited uses and decision owner.	_____
2 Collect information	Explain collection; allow corrections; minimize data; protect privacy.	_____
3 Review a recommendation	Check missing context, incomplete records and uneven outcomes.	_____
4 Before the decision	Require informed human review of all applicants and relevant context.	_____
5 Communicate the outcome	Explain the outcome and offer an accessible correction and appeal pathway.	_____
6 After the process	Monitor results, gather feedback, document concerns and revise or stop use.	_____

MAKE HUMAN OVERSIGHT OPERATIONAL

Name who can pause use, what evidence they inspect, how a student can correct a record, and who documents the final decision. A human signature alone is not meaningful review.

TEACHER GUIDE

Discussion Protocol

Structured dialogue for evidence, counterclaims, and revision

PROTOCOL: CLAIM - EVIDENCE - SAFEGUARD - CHALLENGE

1. Claim:

Each team states whether the current recommendation should be accepted, paused, or rejected.

2. Evidence:

The team cites at least three evidence cards and one stakeholder perspective.

3. Safeguard:

The team proposes at least three changes to the decision process.

4. Challenge:

Another team asks one question that exposes a tradeoff, missing detail, or implementation concern.

5. Revision:

The presenting team revises or defends its recommendation.

Facilitator Questions

- Does efficiency justify using the system in this way?
- Would the process still be acceptable if the result affected you?
- Which information should never be used?
- How might a recommendation influence a reviewer even when humans make the final decision?
- What does a meaningful explanation look like?
- Who should be responsible when staff members rely on a system they do not fully understand?
- What evidence would show that the process is working fairly over time?
- When should the school stop using the system?

Common Misconceptions to Surface

"The system only recommends, so it cannot cause harm."

Recommendations can shape attention, expectations, and access.

"More data automatically creates a fairer result."

Data can still be incomplete, irrelevant, inaccurate, or unevenly distributed.

"A human made the final decision, so the process is accountable."

Human review must be informed, meaningful, and documented.

"Students can appeal, so the system is fair."

Appeals must be accessible, timely, understandable, and capable of changing an outcome.

ASSESSMENT

Performance Rubric

Evaluate reasoning, evidence, safeguards, and communication

Criterion	4 - Advanced	3 - Proficient	2 - Developing	1 - Beginning
Understanding of the decision process	Explains purpose, inputs, recommendation, and human role accurately.	Explains most elements with minor gaps.	Shows partial understanding or confuses recommendation with final decision.	Provides limited or inaccurate explanation.
Evidence analysis	Uses multiple relevant evidence cards and identifies missing or uneven information.	Uses relevant evidence with some analysis.	Uses limited evidence or summarizes without analysis.	Makes unsupported claims.
Fairness and stakeholder reasoning	Examines impact across multiple stakeholders and identifies meaningful fairness concerns.	Considers more than one stakeholder and a plausible fairness concern.	Mentions impact but provides limited perspective-taking.	Does not consider affected stakeholders.
Safeguard design	Proposes specific, feasible safeguards across the decision process.	Proposes useful safeguards with some implementation detail.	Proposes general safeguards with limited detail.	Safeguards are missing or unrelated.
Accountability and human oversight	Clearly identifies who reviews, decides, explains, and remains accountable.	Identifies a human decision-maker and review role.	Mentions human review without clear responsibility.	Leaves final authority unclear.
Communication and defense	Presents a coherent recommendation, responds to challenge, and revises when appropriate.	Communicates a clear recommendation and responds to questions.	Recommendation is understandable but weakly defended.	Recommendation is unclear or unsupported.

USE SCORES TO GUIDE FEEDBACK

Six criteria, 24 points possible. Use the descriptors to identify strengths and next steps; the total is a classroom guide, not a validated mastery threshold. Review individual reasoning as well as the team product.

STUDENT ASSESSMENT

Exit Ticket and Reflection

Individual evidence of understanding

1. Who should be accountable for the final internship decision? Explain why.

2. Identify one piece of missing or incomplete information that could affect Jordan's recommendation.

3. Which safeguard would most improve the process? Explain how it would help.

4. Complete the sentence: AI can support a high-impact decision when...

5. Complete the sentence: AI should not be used to...

OPTIONAL CONFIDENCE CHECK

Circle one response for each statement: 1 = Not yet confident 2 = Somewhat confident 3 = Confident 4 = Very confident

I can distinguish an AI recommendation from a final human decision.

☐ 1 ☐ 2 ☐ 3 ☐ 4

I can identify missing context and possible fairness concerns.

☐ 1 ☐ 2 ☐ 3 ☐ 4

I can explain why transparency and appeal rights matter.

☐ 1 ☐ 2 ☐ 3 ☐ 4

I can design safeguards for responsible AI-supported decisions.

☐ 1 ☐ 2 ☐ 3 ☐ 4

EXTENSION OPTIONS

Adapt the Lesson Across High School Subjects

Flexible pathways for deeper learning

English / Media Literacy

Write an argument evaluating whether the school's process is responsible. Include a claim, counterclaim, evidence, and proposed safeguards.

Social Studies / Civics

Compare the decision process with principles of due process, transparency, public accountability, and equal access.

Computer Science

Create a simplified decision model and test how changing weights or missing values affects recommendations.

Mathematics / Data Science

Analyze a fictional dataset for patterns, missing values, unequal representation, and outcome differences.

CTE / Career Readiness

Redesign the internship selection process to recognize formal coursework, transferable skills, work experience, and community-based learning.

Business / Entrepreneurship

Evaluate the risks and benefits of adopting an AI-supported screening tool. Prepare an implementation recommendation for leadership.

Advisory / Student Leadership

Develop a student-facing bill of rights for AI-supported school decisions.

Art / Digital Media

Create an infographic that explains the five-part Responsible AI Decision Framework.

FAMILY AND COMMUNITY CONNECTION

AI-Supported Decisions Beyond School

Extend the conversation to everyday life and future careers

FAMILY DISCUSSION PROMPT

Artificial intelligence can support decisions in education, employment, finance, health care, transportation, media, and public services. Ask a family or community member to discuss a decision they believe should always include meaningful human review.

Conversation Questions

- 1 Where might an AI-supported recommendation affect a person's opportunity or access?
- 2 What information should people be allowed to review or correct?
- 3 What type of explanation should an organization provide?
- 4 Who should be responsible when an automated recommendation is wrong?
- 5 What makes an appeal process meaningful rather than symbolic?

Student Reflection

Summarize one insight from the conversation and connect it to the lesson's five-part framework: Purpose, Data, Fairness, Transparency, and Human Oversight.

PRIVACY REMINDER

Students should not share private medical, financial, disciplinary, immigration, or other sensitive personal information during this activity.

IMPLEMENTATION NOTES

Adaptation, Accessibility, and Responsible Use

Guidance for educators and school leaders

Differentiation and Accessibility

- Provide the case study in print, digital, audio, or teacher-read formats.
- Reduce the number of evidence cards or pre-highlight key details for students who need additional support.
- Use sentence starters for claims, evidence, counterclaims, and safeguards.
- Assign stakeholder roles strategically to support participation and perspective-taking.
- Allow students to respond through writing, discussion, graphic organizers, or recorded presentations.
- Define key terms before the lesson: recommendation, data, bias, transparency, oversight, accountability, and appeal.
- Avoid requiring students to disclose personal experiences with school discipline, disability, immigration, financial hardship, or other sensitive contexts.

Responsible Technology Use

- No student AI account is required for this lesson.
- Do not upload identifiable student information to public AI tools.
- When using an AI tool for extension work, provide an approved prompt and require students to verify outputs.
- Treat generated content as a draft or source for evaluation, not as an authority.
- Require disclosure when AI tools contribute to a student product.
- Follow district-approved technology, privacy, accessibility, and academic-integrity expectations.

Evidence of Learning

- Completed Responsible AI Decision Framework
- Team-designed decision process
- Oral defense and response to challenge
- Individual exit ticket
- Optional family/community reflection
- Teacher observation using the performance rubric

Consulting Use

This flagship lesson can anchor a short pilot, educator professional-learning session, curriculum review, or district AI-readiness conversation. It is intentionally adaptable across subject areas and technology environments.

Explicit limits and stop authority

For any tool-based extension, name the permitted inputs and actions, the output reviewer, and the person who can stop use. Use fictional records only; no ranking or screening of actual students.

MEETING HANDOUT

One-Page Lesson Summary

Use during a district or school consultation conversation

WHY THIS LESSON MATTERS

High school students are already encountering AI-supported recommendations and generated content. This lesson helps them move beyond tool use toward judgment, evidence analysis, fairness, transparency, and accountability.

What Students Do

- Investigate a fictional internship recommendation
- Analyze evidence and stakeholder perspectives
- Identify missing context and potential unfairness
- Design safeguards and an appeal pathway
- Defend a recommendation using evidence
- Explain where human judgment and accountability belong

What Educators Can Observe

- Quality of student reasoning
- Ability to distinguish recommendation from decision
- Use of evidence and counterclaims
- Understanding of privacy, fairness, and transparency
- Ability to propose practical safeguards
- Readiness to apply the framework to other contexts

Possible Pilot Uses

SINGLE LESSON

Advisory, social studies, English, computer science, business, or CTE

MINI-UNIT

Three lessons on decisions, information quality, and accountability

PROFESSIONAL LEARNING

Teachers experience the lesson as learners, then plan subject-specific adaptations

STUDENT LEADERSHIP

Student advisory group reviews school AI expectations and communication

CORE MESSAGE

AI may support a decision. People must still define the purpose, examine the evidence, explain the outcome, provide a way to challenge it, and remain accountable.